

Machine Learning

Volker Roth

Department of Mathematics & Computer Science
University of Basel

Section 9

Mixture Models

Structure and mixtures

- Assume that input examples come in different potentially **unobserved types (groups, clusters, etc.)**.
- Assume that
 - ① there are m underlying types $z = 1, \dots, m$;
 - ② each type z occurs with probability $P(z)$;
 - ③ examples of type z distributed according to $p(\mathbf{x}|z)$.
- According to this model, each observed \mathbf{x} comes from a **mixture distribution:**

$$p(\mathbf{x}) = \sum_{j=1}^m \underbrace{P(z=j)}_{\pi_j} p(\mathbf{x}|z=j, \theta_j)$$

- In many practical data analysis problems (such as probabilistic clustering), we want to **estimate** such parametric models from **samples** $\{\mathbf{x}_1, \dots, \mathbf{x}_n\}$. In particular, we are often interested in finding the types that have generated the examples.

Mixture of Gaussians

A mixture of Gaussians model has the form

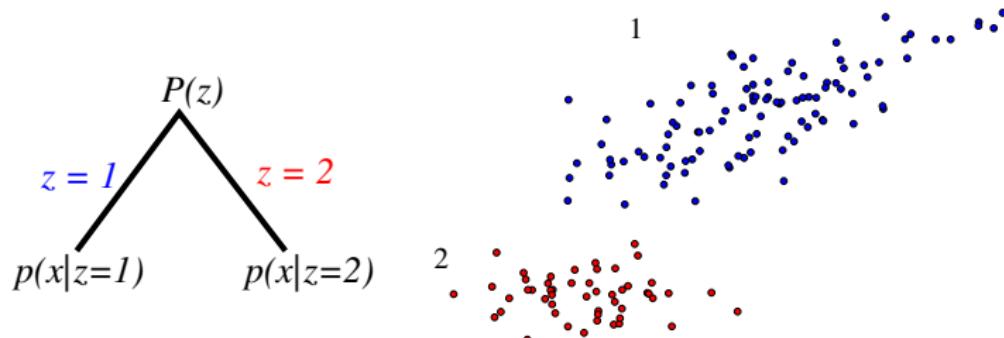
$$p(\mathbf{x}|\boldsymbol{\theta}) = \sum_{j=1}^m \pi_j \mathcal{N}(\mathbf{x}|\boldsymbol{\mu}_j, \boldsymbol{\Sigma}_j),$$

where $\boldsymbol{\theta} = \pi_1, \dots, \pi_m, \boldsymbol{\mu}_1, \dots, \boldsymbol{\mu}_m, \boldsymbol{\Sigma}_1, \dots, \boldsymbol{\Sigma}_m$ contains all the parameters. $\{\pi_j\}$ are the **mixing proportions**.



Mixture densities

- Data generation process:

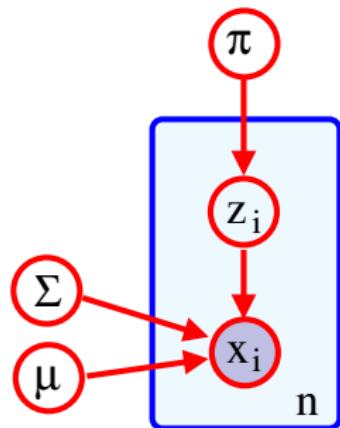


$$p(\mathbf{x}|\boldsymbol{\theta}) = \sum_{j=1}^m \pi_j p(\mathbf{x}|\boldsymbol{\mu}_j, \boldsymbol{\Sigma}_j)$$

- Any data point \mathbf{x} could have been **generated in two ways**.
~~ the responsible component needs to be **inferred**.

Mixtures as Latent Variable Models

- In the model $p(\mathbf{x}|z = j, \theta)$ the class indicator variable z is **latent**. This is an example of a large class of **latent variable models** (LVM).
- Bayesian network (DAG) = graphical representation of the joint distribution of RVs (nodes) as $P(x_1, \dots, x_n) = \prod_{i=1}^n P(x_i|parents(x_i))$



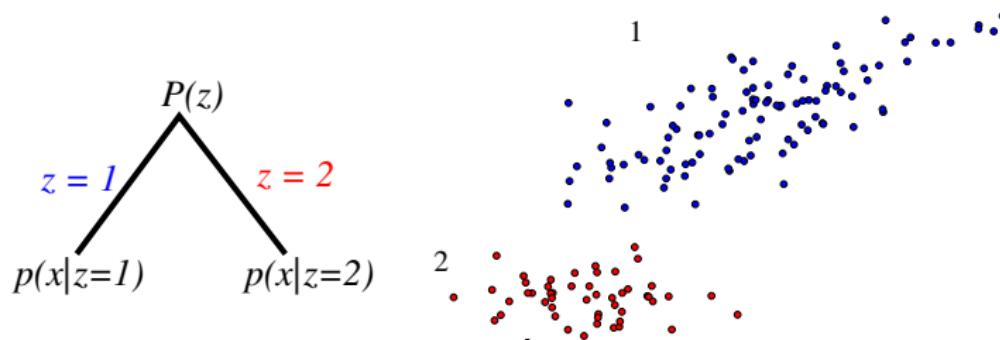
$$\begin{aligned} p(\mathbf{x}_i|\theta) &= \sum_{z_i} p(\mathbf{x}_i, z_i|\theta) \\ &= \sum_{z_i} p(\mathbf{x}_i|\mu, \Sigma, z_i)p(z_i|\pi). \end{aligned}$$

Mixture densities

- Consider a two component mixture of Gaussians model.

$$p(\mathbf{x}|\boldsymbol{\theta}) = \pi_1 p(\mathbf{x}|\boldsymbol{\mu}_1, \boldsymbol{\Sigma}_1) + \pi_2 p(\mathbf{x}|\boldsymbol{\mu}_2, \boldsymbol{\Sigma}_2)$$

- If we knew the generating component $z_i = \{1, 2\}$ for each example \mathbf{x}_i , then the estimation would be easy.



- In particular, we can estimate each Gaussian independently.

Mixture density estimation

- Let $\delta(j|i)$ be an indicator function of whether example i is labeled j .
Then for each $j = 1, 2$

$$\hat{\pi}_j \leftarrow \frac{\hat{n}_j}{n}, \quad \text{where } \hat{n}_j = \sum_{i=1}^n \delta(j|i)$$

$$\hat{\mu}_j \leftarrow \frac{1}{\hat{n}_j} \sum_{i=1}^n \delta(j|i) \mathbf{x}_i$$

$$\hat{\Sigma}_j \leftarrow \frac{1}{\hat{n}_j} \sum_{i=1}^n \delta(j|i) (\mathbf{x}_i - \hat{\mu}_j)(\mathbf{x}_i - \hat{\mu}_j)^t$$

Mixture density estimation

- We don't have such labels... but we can guess what the labels might be based on our current distribution.
- One possible choice: evaluate posterior probability that an observed \mathbf{x} was generated from first component

$$\begin{aligned} P(z=1|\mathbf{x}, \theta) &= \frac{P(z=1) \cdot p(\mathbf{x}|z=1)}{\sum_{j=1,2} P(z=j) \cdot p(\mathbf{x}|z=j)} \\ &= \frac{\pi_1 p(\mathbf{x}|\boldsymbol{\mu}_1, \Sigma_1)}{\sum_{j=1,2} \pi_j p(\mathbf{x}|\boldsymbol{\mu}_j, \Sigma_j)} \end{aligned}$$

↔ Information about the component responsible for generating \mathbf{x} .

- **Soft labels or posterior probabilities**

$$\hat{p}(j|i) \leftarrow P(z_i = j | \mathbf{x}_i, \theta),$$

where $\sum_{j=1,2} \hat{p}(j|i) = 1, \forall i = 1, \dots, n$.

The EM algorithm: iteration k

- **E-step:** softly assign examples to mixture components

$$\hat{p}(j|i) \leftarrow P(z_i = j | \mathbf{x}_i, \boldsymbol{\theta}^t), \forall j = 1, 2 \text{ and } i = 1, \dots, n.$$

Note: superscript is time index.

- **M-step:** estimate new mixture parameters $\boldsymbol{\theta}^{t+1}$ based on the soft assignments (can be done separately for the two Gaussians)

$$\hat{\pi}_j \leftarrow \frac{\hat{n}_j}{n}, \quad \text{where } \hat{n}_j = \sum_{i=1}^n \hat{p}(j|i)$$

$$\hat{\mu}_j \leftarrow \frac{1}{\hat{n}_j} \sum_{i=1}^n \hat{p}(j|i) \mathbf{x}_i$$

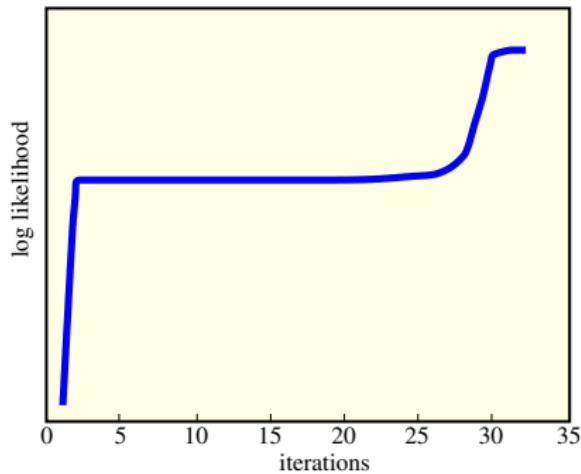
$$\hat{\Sigma}_j \leftarrow \frac{1}{\hat{n}_j} \sum_{i=1}^n \hat{p}(j|i) (\mathbf{x}_i - \hat{\mu}_j)(\mathbf{x}_i - \hat{\mu}_j)^t$$

The EM-algorithm: Convergence

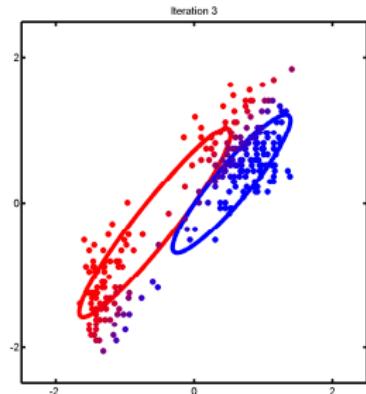
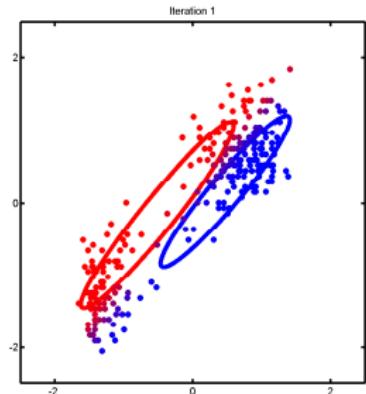
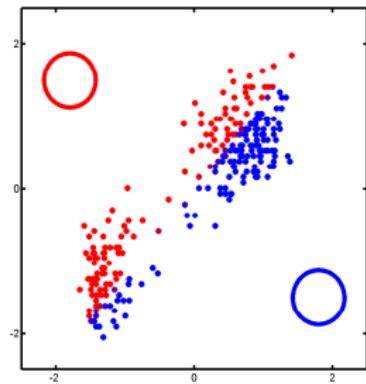
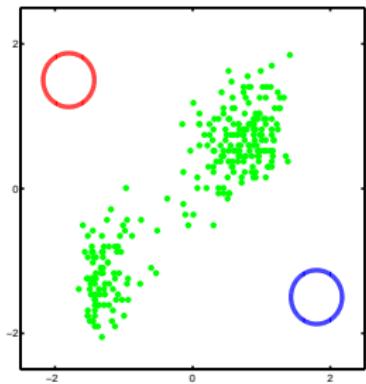
The EM-algorithm **monotonically increases the log-likelihood** of the training data (we will show this later). In other words,

$$l(\theta^0) < l(\theta^1) < l(\theta^2) < \dots \text{ until convergence}$$

$$l(\theta^t) = \sum_{i=1}^n \log p(x_i | \theta^t).$$



Mixture density estimation: example



Mixture density estimation: example

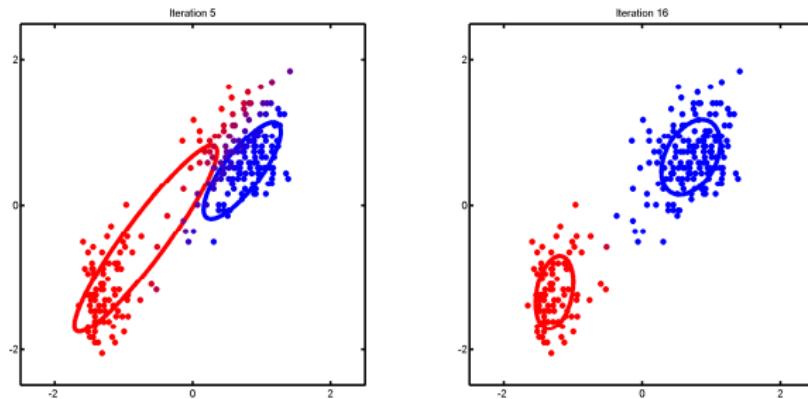
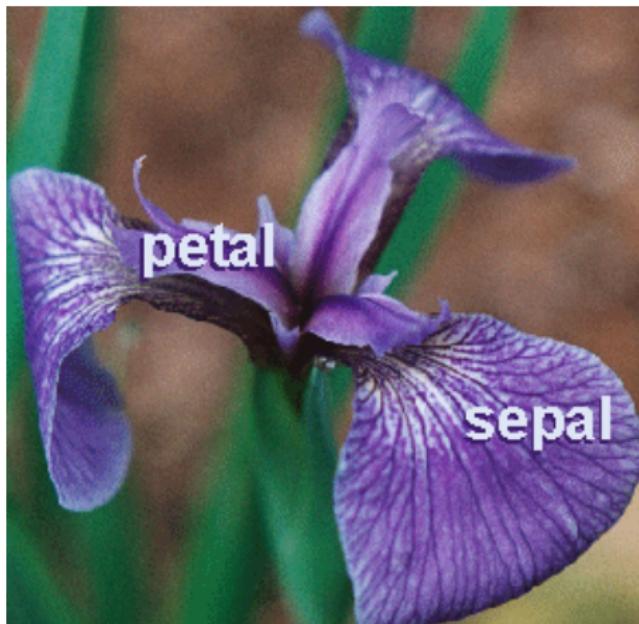


Fig. 11.11 in K. Murphy

EM example: Iris data

- The famous (Fisher's or Anderson's) iris data set gives the measurements in centimeters of the variables **sepal length and width** and **petal length and width**, respectively, for 50 flowers from each of **3 species of iris**.
- The species are **Iris setosa, versicolor, and virginica**.



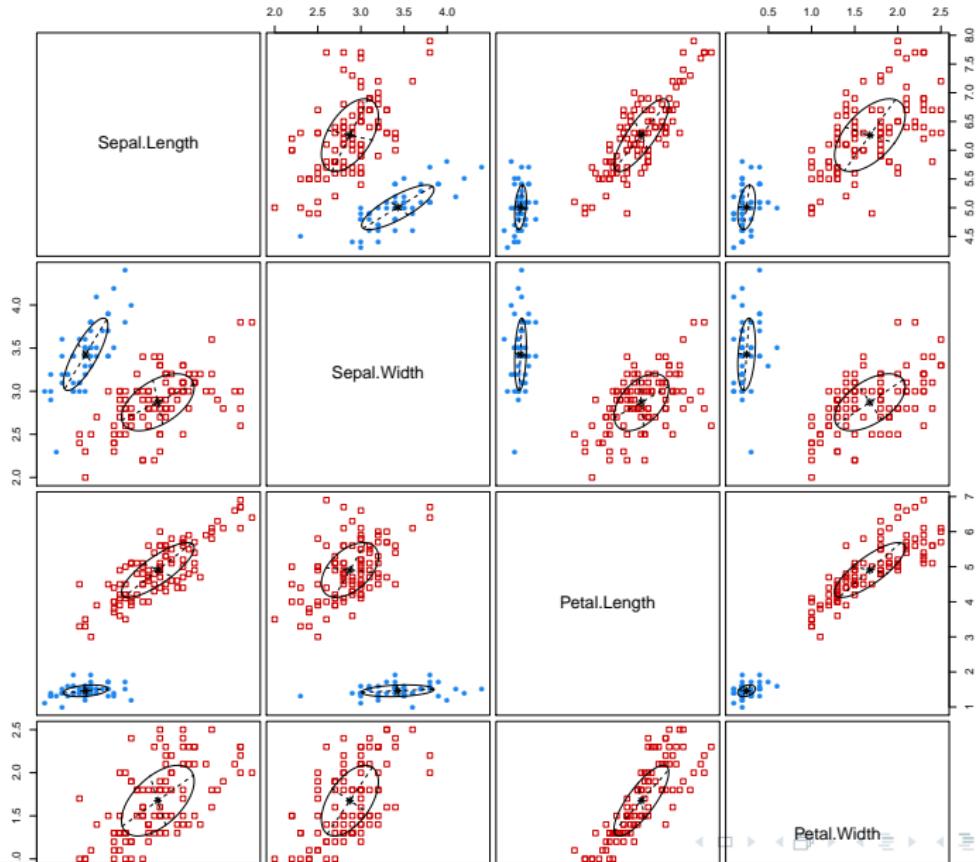
Bayesian model selection for mixture models

As a simple strategy for selecting the appropriate number of mixture components, we can find m that minimizes the overall description length (cf. BIC):

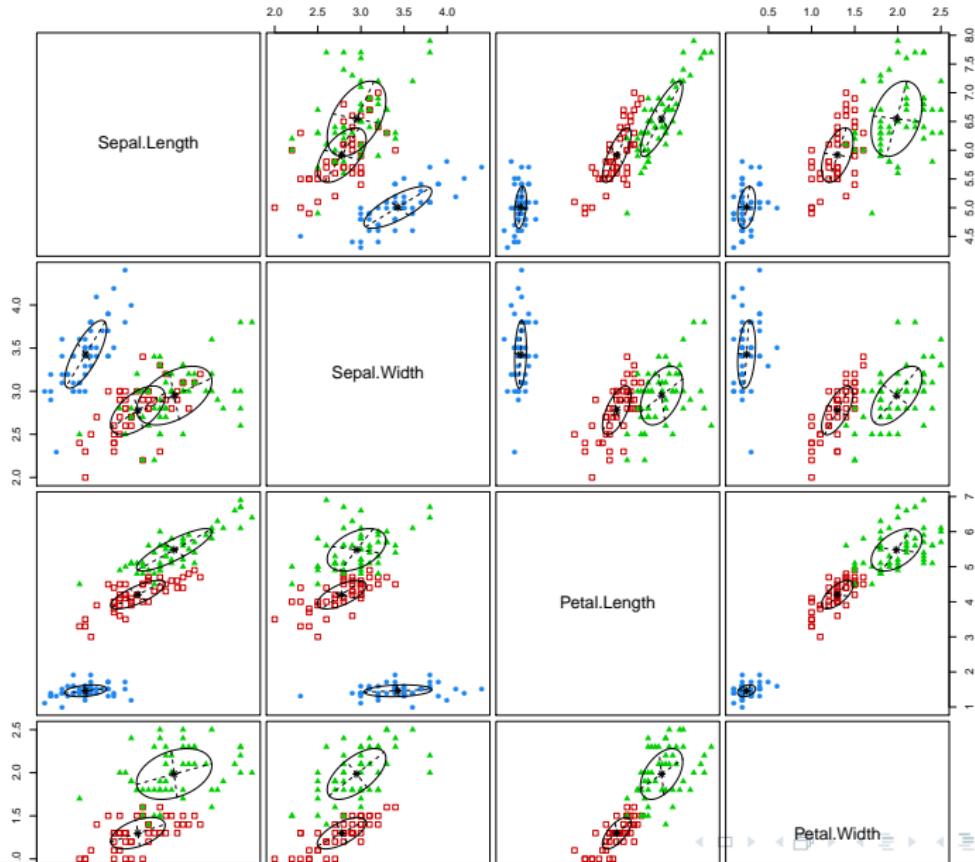
$$DL \approx -\log p(\text{data}|\hat{\theta}_m) + \frac{d_m}{n} \log(n)$$

- n is the number of training points,
- $\hat{\theta}_m$ are the maximum likelihood parameters for the m -component mixture, and
- d_m is the number of parameters in the m -component mixture.

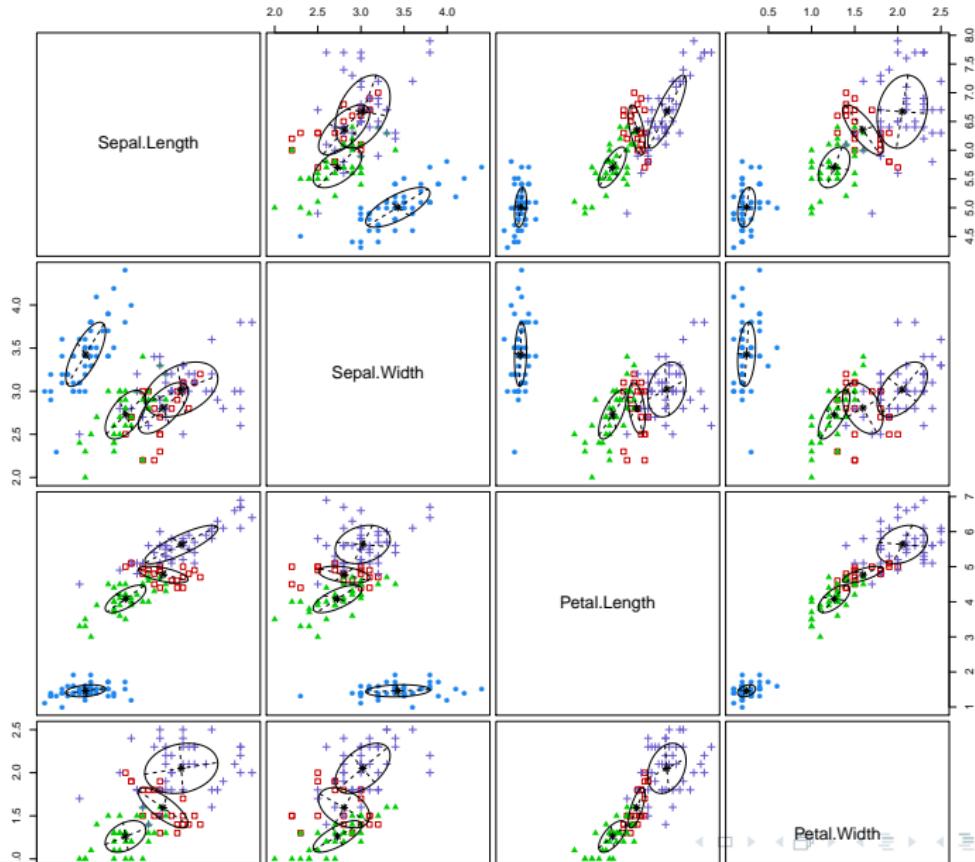
Model selection example: Iris data, $m = 2$



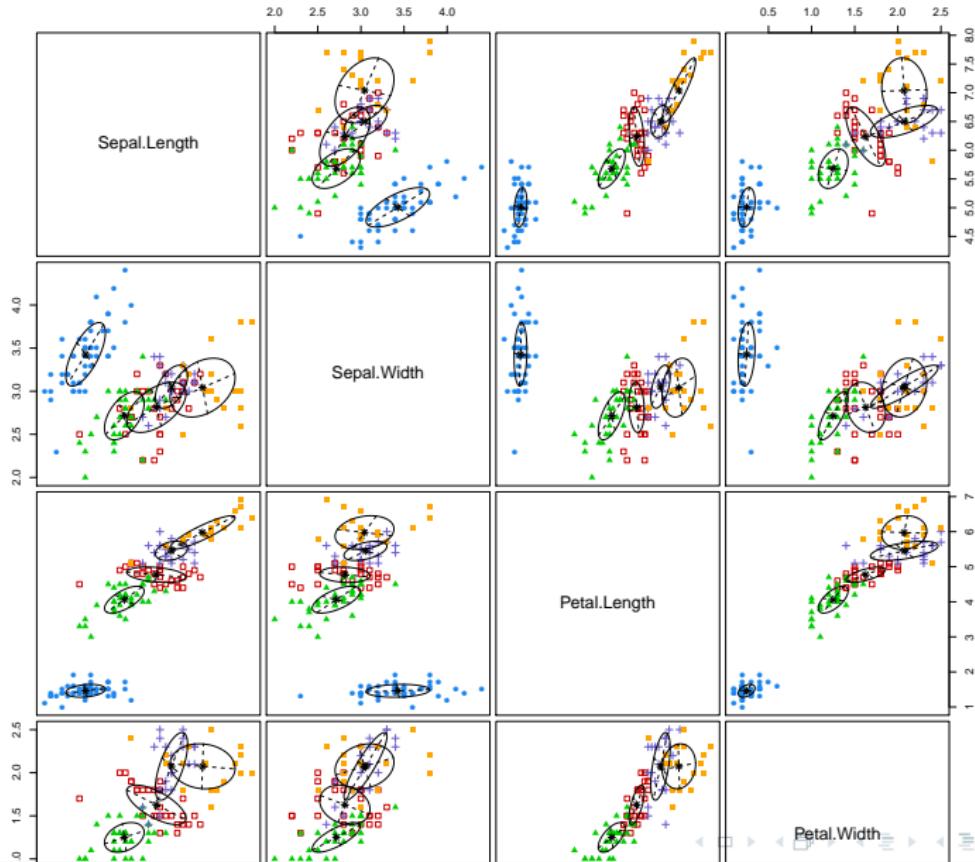
Model selection example: Iris data, $m = 3$



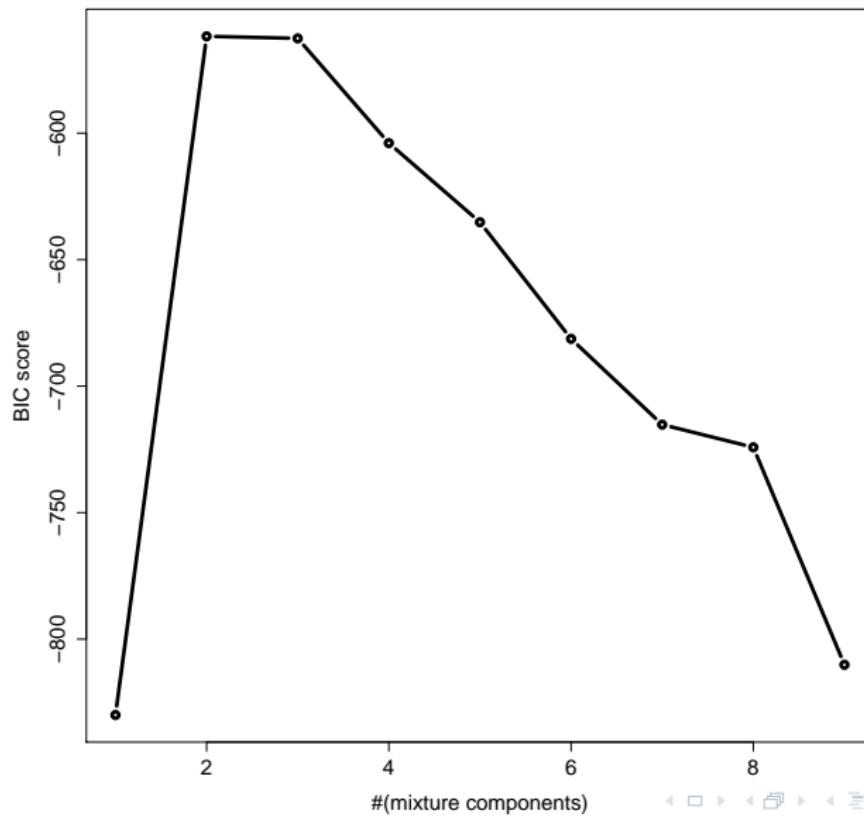
Model selection example: Iris data, $m = 4$



Model selection example: Iris data, $m = 5$



Model selection example: Iris data, BIC



The EM-algorithm: Convergence

Step 0: specify the initial setting of the parameters $\theta = \theta^0$.

E-step: complete the incomplete data (missing z) with the posterior probabilities (“soft labels”)

$$P(z = j | \mathbf{x}_i, \theta^t), \quad j = 1, \dots, m, \quad i = 1, \dots, n.$$

M-step: find the new setting of the parameters θ^{t+1} by maximizing the log-likelihood of the inferred (or “expected complete”) data

$$\theta^{t+1} = \arg \max_{\theta} \underbrace{\sum_{i=1}^n \sum_{j=1}^m P(z = j | \mathbf{x}_i, \theta^t) \log[p_j p(\mathbf{x}_i | \theta_j)]}_{\text{inferred (= expected complete) log-likelihood } Q(\theta, \theta^t)} .$$

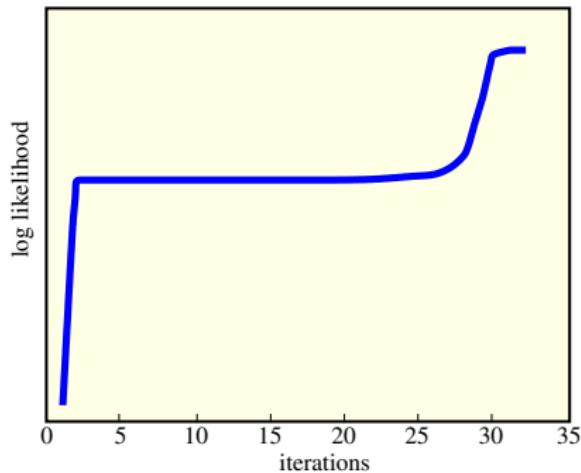
The expected complete log-likelihood $Q(\theta, \theta^t)$ is called
auxiliary objective.

The EM-algorithm: Convergence

The EM-algorithm **monotonically increases the log-likelihood** of the training data. In other words,

$$l(\theta^0) < l(\theta^1) < l(\theta^2) < \dots \text{ until convergence}$$

$$l(\theta^t) = \sum_{i=1}^n \log p(x_i | \theta^t).$$



Jensen's inequality

- Convex function: secant line above graph of the function
~~> Jensen's inequality for two points.
- Secant line consists of weighted means of the convex function.
For $a \in [0, 1]$:

$$af(x_1) + (1 - a)f(x_2).$$

Graph: convex function of the weighted means:

$$f(ax_1 + (1 - a)x_2).$$

- Thus, Jensen's inequality is

$$f(ax_1 + (1 - a)x_2) \leq af(x_1) + (1 - a)f(x_2).$$

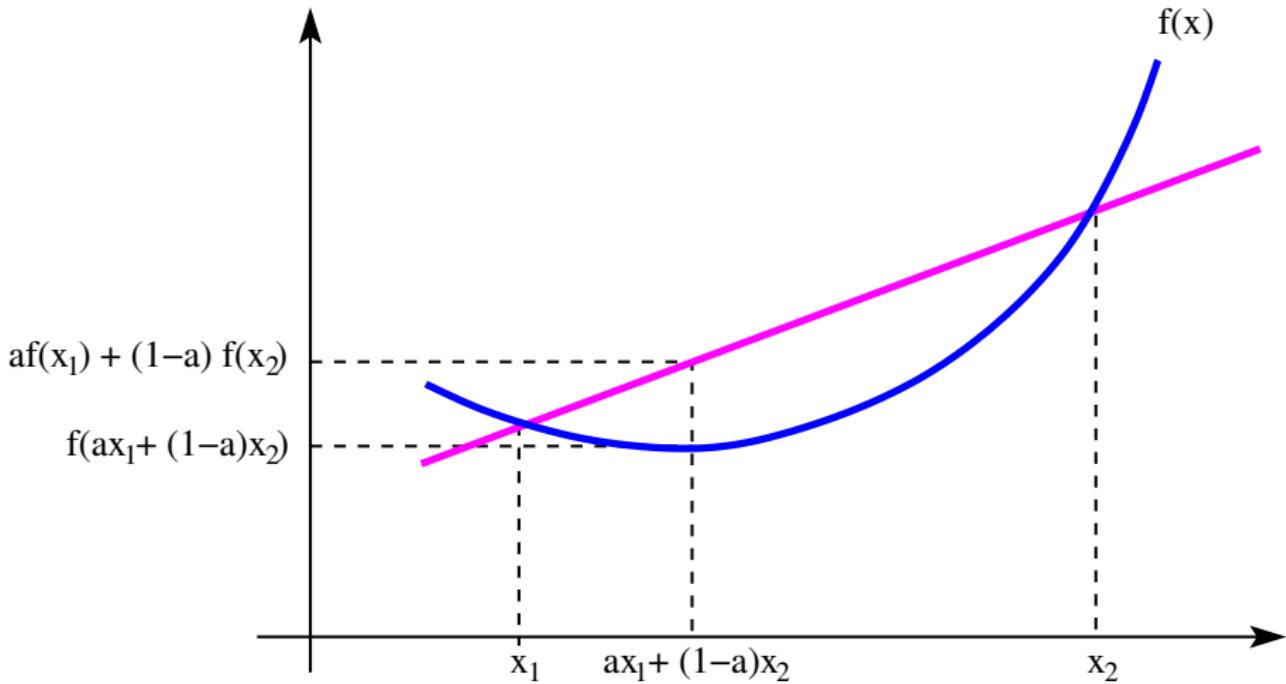
- Probability theory: if X is a RV and φ a convex function, then

$$\varphi(E[X]) \leq E[\varphi(X)].$$

- φ convex ~> $\psi := -\varphi$ concave:

$$\psi(E[X]) \geq E[\psi(X)]. \quad \text{Example:} \quad \log(E[X]) \geq E[\log(X)].$$

Jensen's inequality



Non-negativity of KL divergence

$$\begin{aligned}-\mathbb{KL}(p(x) \| q(x)) &= \int p(x) \log \left(\frac{q(x)}{p(x)} \right) dx \\ (\text{Jensen's inequality}) \quad &\leq \log \left(\int p(x) \frac{q(x)}{p(x)} dx \right) \\ &= \log \left(\int q(x) dx \right) \\ &= \log(1) = 0\end{aligned}$$

This is also called **Gibbs' inequality.**

The EM-algorithm: Theoretical basis

Consider distribution $q(z_i)$ over latent assignment variables.

Log-likelihood:

$$\begin{aligned} I(\theta) &= \sum_{i=1}^n \log p(x_i | \theta) \\ &= \sum_{i=1}^n \log \sum_{z_i} p(x_i, z_i | \theta) \\ &= \sum_{i=1}^n \log \sum_{z_i} q(z_i) \frac{p(x_i, z_i | \theta)}{q(z_i)} \\ &= \sum_{i=1}^n \log \mathbb{E}_{q_i} \frac{p(x_i, z_i | \theta)}{q(z_i)} \\ (\text{Jensen's inequality}) \quad &\geq \sum_{i=1}^n \sum_{z_i} q(z_i) \log \frac{p(x_i, z_i | \theta)}{q(z_i)} \\ &=: Q(\theta, q). \end{aligned}$$

- Lower bound

$$I(\theta) \geq Q(\theta, q) := \sum_{i=1}^n \sum_{z_i} q(z_i) \log \frac{p(x_i, z_i | \theta)}{q(z_i)}$$

valid for any positive distribution q . Which one should we choose?

- Intuition: pick the q that yields the tightest lower bound.
This will be the **E-step**.
- At time t , assume we have chosen q^t based on current parameters θ^t .
In the next **M-step** we maximize the
expected complete log-likelihood:

$$\theta^{t+1} = \arg \max_{\theta} Q(\theta, \theta^t) = \arg \max_{\theta} \sum_{i=1}^n \mathbb{E}_{q_i^t} \log p(x_i, z_i | \theta)$$

- Last equation follows from

$$Q(\theta, q) = \underbrace{\sum_{i=1}^n \mathbb{E}_{q_i} \log p(x_i, z_i | \theta)}_{\text{Expected complete log-l}} + \sum_{i=1}^n \underbrace{\left[- \sum_{z_i} q(z_i) \log q(z_i) \right]}_{h(q_i), \text{ independent of } \theta}.$$

The E-step

Re-write lower bound as

$$Q(\theta, q) = \sum_i L(\theta, q_i),$$

with

$$\begin{aligned}L(\theta, q_i) &= \sum_{z_i} q(z_i) \log \frac{p(\mathbf{x}_i, z_i | \theta)}{q(z_i)} \\&= \sum_{z_i} q(z_i) \log \frac{p(z_i | \mathbf{x}_i, \theta) p(\mathbf{x}_i | \theta)}{q(z_i)} \\&= \sum_{z_i} q(z_i) \log \frac{p(z_i | \mathbf{x}_i, \theta)}{q(z_i)} + \sum_{z_i} q(z_i) \log p(\mathbf{x}_i | \theta) \\&= \underbrace{-\mathbb{KL}(q(z_i) \| p(z_i | \mathbf{x}_i, \theta))}_{\text{always } \geq 0, \text{ and } = 0, \text{ if } q=p} + \underbrace{\log p(\mathbf{x}_i | \theta)}_{\text{independent of } q_i}.\end{aligned}$$

The E step

For $q_i^t(z_i) = p(z_i|\mathbf{x}_i, \theta^t)$, the KL divergence is zero, and $L(\theta^t, q_i)$ is maximized over all possible distributions q_i :

$$q_i^t(z_i) = p(z_i|\mathbf{x}_i, \theta^t) = \arg \max_{q_i} L(\theta^t, q_i) \quad (\rightsquigarrow \text{E-step})$$

$$L(\theta^t, q_i^t) = \log p(\mathbf{x}_i|\theta^t)$$

$$Q(\theta^t, \theta^t) = \sum_i \log p(\mathbf{x}_i|\theta^t) = l(\theta^t)$$

- ~ lower bound “touches” the log-likelihood
- ~ after the E-step, the auxiliary objective equals the log-likelihood
- ~ lower bound is tight after the E-step.

The E step

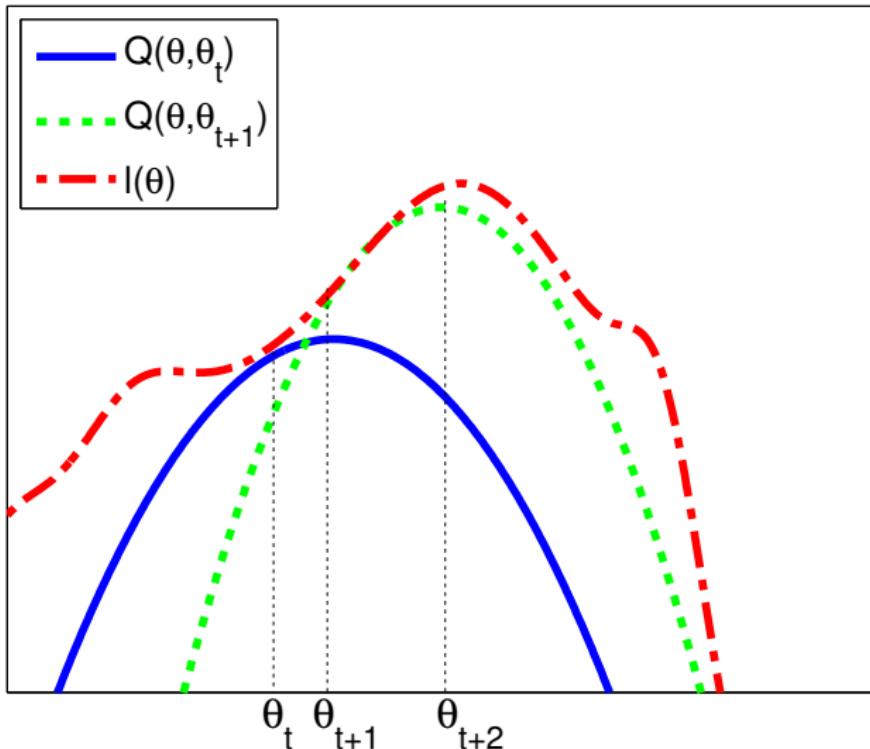


Fig 11.16 in K. Murphy

EM-algorithm: max-max and monotonicity

We can now rewrite the EM-algorithm in terms of two maximization steps involving the auxiliary objective:

E-step: $q^t = \arg \max_q Q(\theta^t, q)$

M-step: $\theta^{t+1} = \arg \max_{\theta} Q(\theta, \theta^t)$.

The monotonic increase of the log-likelihood now follows from

$$\underbrace{I(\theta^{t+1}) \geq Q(\theta^{t+1}, \theta^t)}_{Q(\theta^{t+1}, \bullet) \text{ is lower bound on } I(\theta^{t+1})} \geq Q(\theta^t, \theta^t) = I(\theta^t).$$

Second inequality: $Q(\theta^{t+1}, \theta^t) = \max_{\theta} Q(\theta, \theta^t) \geq Q(\theta^t, \theta^t)$.

Conditional mixtures

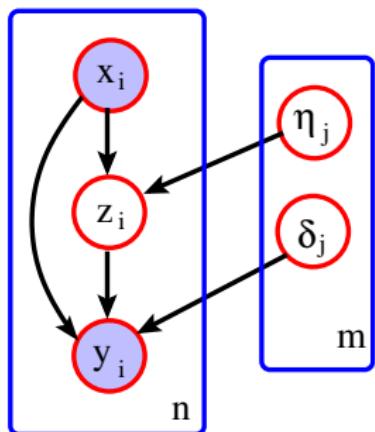
- Some regression or classification problems can be decomposed into easier sub-problems.
- Examples:
 - ▶ style in handwritten character recognition
 - ▶ dialect/accent in speech recognition, etc.
- Each sub-problem could be solved by a specific “expert”.
- The selection of which expert to rely on now depends on the position x in the input space. **Mixtures of experts models.**

Experts (regression)

- Suppose we have several regression “experts” generating conditional Gaussian outputs

$$p(y|\mathbf{x}, z = j, \boldsymbol{\theta}) = \mathcal{N}(y | \boldsymbol{\beta}_j^t \mathbf{x}, \sigma_j^2)$$

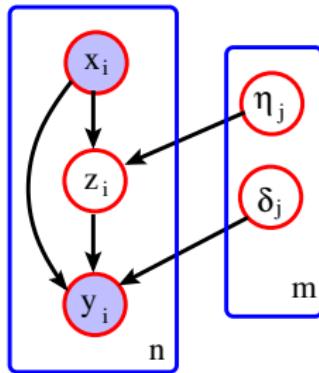
- $\boldsymbol{\delta}_j = \{\boldsymbol{\beta}_j, \sigma_j^2\}$: Parameters of j -th expert.
- Need to find a way of allocating tasks to these experts.
- Parameter vector $\boldsymbol{\theta}$ contains the means and variances of the m experts and the additional parameters $\boldsymbol{\eta}$ of this allocation mechanism:
$$\boldsymbol{\theta} = \{\boldsymbol{\delta}_j, \boldsymbol{\eta}_j\}_{j=1}^m.$$



Joint distribution

From the DAG we conclude:

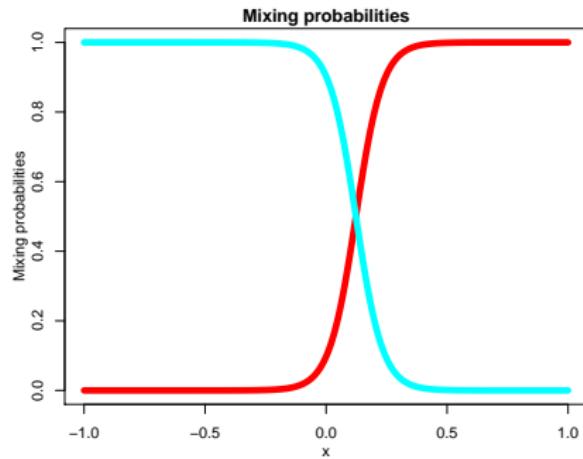
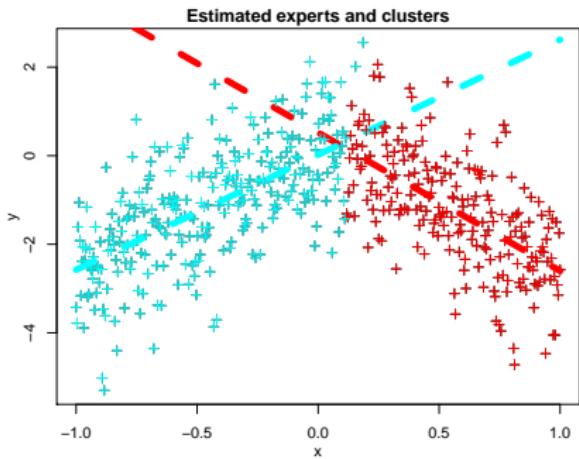
$$\begin{aligned} p(y, z = j | \mathbf{x}) &= p(y | \boldsymbol{\delta}, z = j, \mathbf{x}) P(z = j | \boldsymbol{\eta}, \mathbf{x}) \\ &= p(y | \boldsymbol{\delta}_j, \mathbf{x}) P(z = j | \boldsymbol{\eta}, \mathbf{x}) \\ &= \mathcal{N}(y | \boldsymbol{\beta}_j^T \mathbf{x}, \sigma_j^2) P(z = j | \boldsymbol{\eta}, \mathbf{x}) \end{aligned}$$



Thus, the overall prediction is

$$\begin{aligned} p(y | \mathbf{x}, \theta) &= \sum_j p(y, z = j | \mathbf{x}, \boldsymbol{\eta}, \boldsymbol{\delta}) \\ &= \sum_j P(z = j | \mathbf{x}, \boldsymbol{\eta}) p(y | \mathbf{x}, \boldsymbol{\delta}_j) \\ &= \sum_j P(z = j | \mathbf{x}, \boldsymbol{\eta}) p(y | \mathbf{x}, \boldsymbol{\beta}_j, \sigma_j^2). \end{aligned}$$

Mixtures of experts



Here we need to switch from one linear regression model to another:
 $p(y|x, z = j, \theta) = \mathcal{N}(y|\beta_j^t x, \sigma_j^2)$. The switch can be probabilistic
↔ probabilistic gating function $P(z|x, \eta)$ (right).

Gating network

- A gating network specifies a distribution over m experts, conditionally on the input \mathbf{x} .
- Example: when there are just two experts the gating network can be a logistic regression model

$$P(z = 1|\mathbf{x}, \boldsymbol{\eta}) = \sigma(\boldsymbol{\eta}^t \mathbf{x}),$$

where $\sigma(z) = (1 + e^{-z})^{-1}$ is the logistic function.

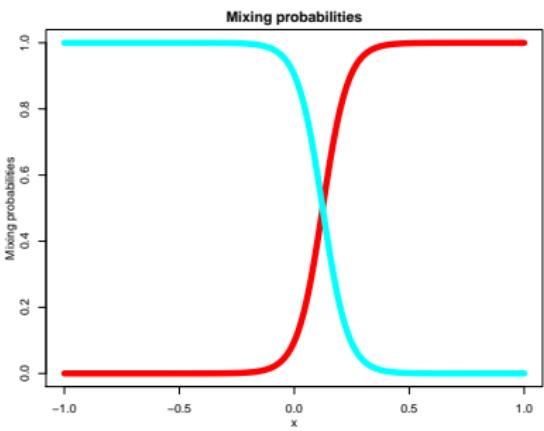
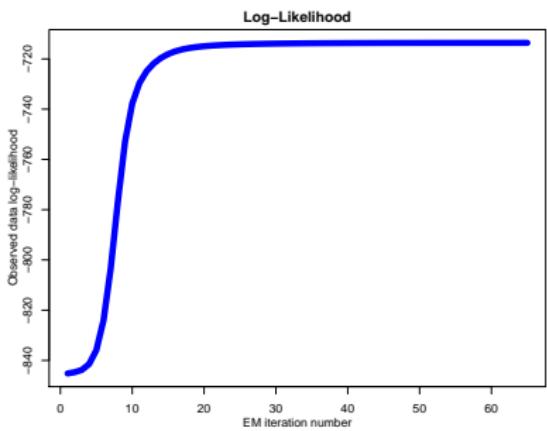
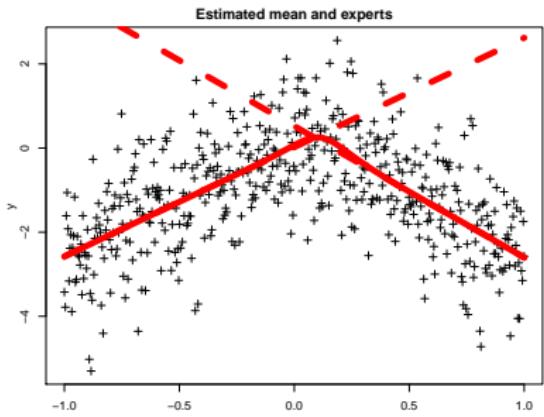
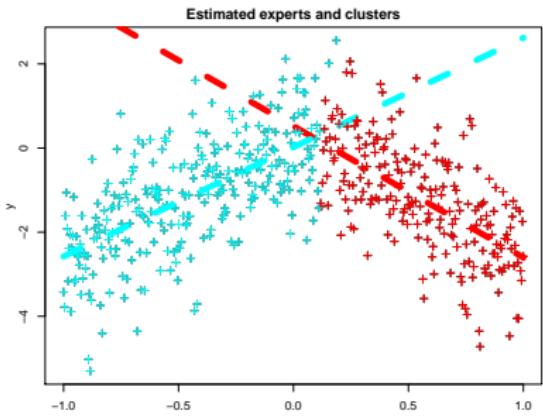
- For $m > 2$, the gating network can be a softmax model

$$P(z = j|\mathbf{x}, \boldsymbol{\eta}) = \frac{\exp(\boldsymbol{\eta}_j^t \mathbf{x})}{\sum_{j'=1}^m \exp(\boldsymbol{\eta}_{j'}^t \mathbf{x})},$$

where $\boldsymbol{\eta} = \{\boldsymbol{\eta}_1, \dots, \boldsymbol{\eta}_m\}$ are the parameters of the gating network.

- Overall prediction

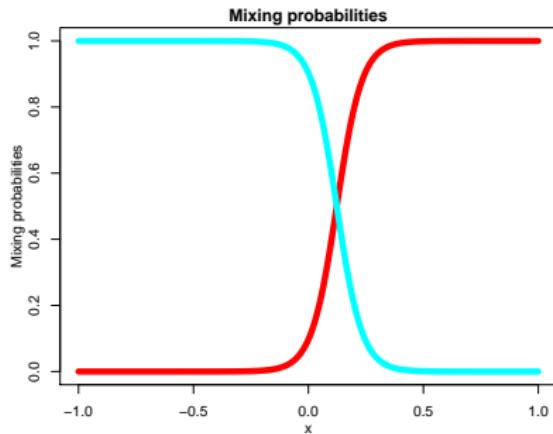
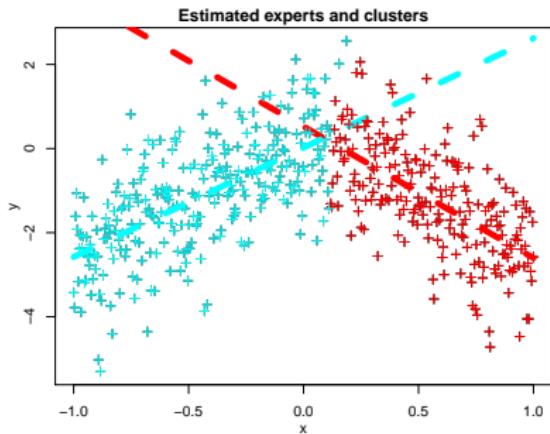
$$p(y|\mathbf{x}, \boldsymbol{\theta}) = \sum_j p(y, z = j|\mathbf{x}, \boldsymbol{\eta}, \boldsymbol{\delta}) = \sum_j P(z = j|\mathbf{x}, \boldsymbol{\eta})p(y|\mathbf{x}, \boldsymbol{\delta}_j).$$



A mixture of experts model: estimation

“Soft labels”: Conditional probability that (\mathbf{x}_i, y_i) came from expert j :

$$\begin{aligned}\hat{P}(j|i) &= P(z = j|\mathbf{x}_i, y_i, \theta) \\ &= \frac{P(z = j|\mathbf{x}_i, \eta^t)p(y_i|\mathbf{x}_i, (\beta_j, \sigma_j^2))}{\sum_{j'=1}^m P(z = j'|\mathbf{x}_i, \eta^t)p(y_i|\mathbf{x}_i, (\beta_{j'}, \sigma_{j'}^2))}\end{aligned}$$



EM for mixtures of experts

E-step: compute soft labels $\hat{P}(j|i)$

M-step: separately re-estimate the experts and the gating network based on these soft assignments:

- ① For each expert j : find $(\hat{\beta}_j, \hat{\sigma}_j^2)$ that maximize

$$\sum_{i=1}^n \hat{P}(j|i) \log p(y_i | \mathbf{x}_i, (\beta_j, \sigma_j^2))$$

~~ linear regression with weighted training set.

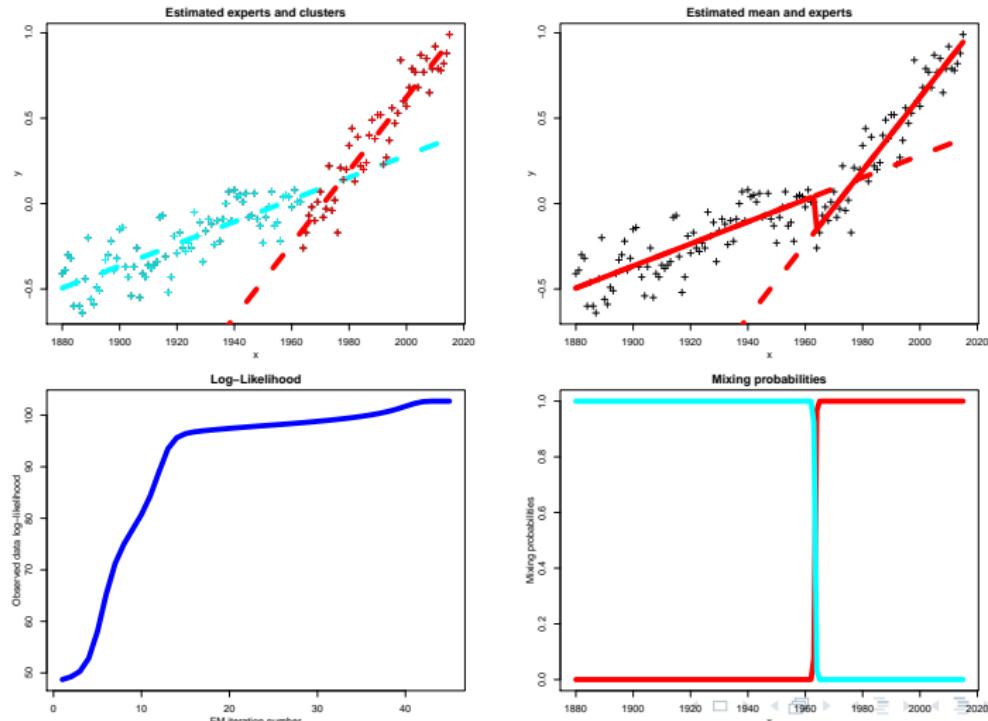
- ② For the gating network: find $\hat{\eta}$ that maximize

$$\sum_{i=1}^n \sum_{j=1}^m \hat{P}(j|i) \log P(j | \mathbf{x}_i, \eta_j)$$

~~ logistic regression with weighted training set.

Real World Example

Global annual temperature anomalies (degrees C) computed using data from land meteorological stations, 1880-2015. Anomalies are relative to the 1951-1980 base period means.



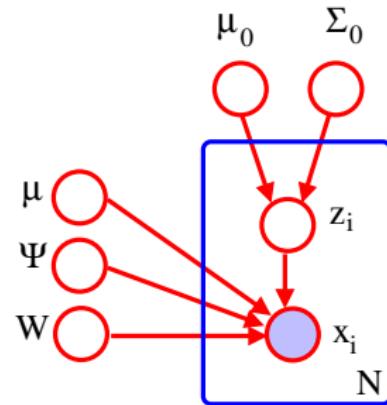
Section 10

Linear latent variable models

Factor analysis

- One problem with mixture models: **only a single latent variable**. Each observation can only come from one of K prototypes.
- Alternative: $\mathbf{z}_i \in \mathbb{R}^k$. Gaussian prior:

$$p(\mathbf{z}_i) = \mathcal{N}(\mathbf{z}_i | \boldsymbol{\mu}_0, \boldsymbol{\Sigma}_0)$$



- For observations $\mathbf{x}_i \in \mathbb{R}^p$, we may use a **Gaussian likelihood**.
- As in linear regression, we assume the mean is a **linear** function:
$$p(\mathbf{x}_i | \mathbf{z}_i, \theta) = \mathcal{N}(W\mathbf{z}_i + \boldsymbol{\mu}, \boldsymbol{\Psi}),$$

W: **factor loading matrix**, and $\boldsymbol{\Psi}$: **covariance matrix**.
- We take $\boldsymbol{\Psi}$ to be **diagonal**, since the whole point of the model is to “force” \mathbf{z}_i to **explain the correlation**.

Factor analysis: generative process

Generative process ($k = 1$, $p = 2$, diagonal Ψ):

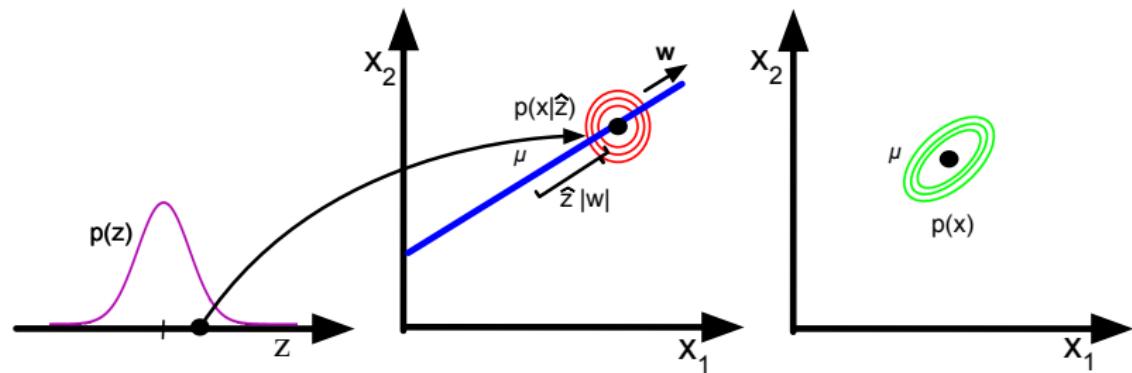


Figure 12.1 in K. Murphy

We take an isotropic Gaussian “spray can” and slide it along the 1d line defined by $wz_i + \mu$. This induces a correlated Gaussian in 2d.

Inference of the latent factors

- We hope that the latent factors z will reveal something interesting about the data \rightsquigarrow compute posterior over the latent variables:

$$p(z_i | \mathbf{x}_i, \theta) = \mathcal{N}(z_i | \mathbf{m}_i, \Sigma)$$

$$\Sigma = (\Sigma_0^{-1} + W^t \Psi^{-1} W)^{-1}$$

$$\mathbf{m}_i = \Sigma_i (W^t \Psi^{-1} (\mathbf{x}_i - \boldsymbol{\mu}) + \Sigma_0^{-1} \boldsymbol{\mu}_0)$$

- The posterior means \mathbf{m}_i are called the **latent scores, or latent factors.**

Example

- Example from (Shalizi 2009). $p = 11$ variables and $n = 387$ cases describing aspects of cars: engine size, #(cylinders), miles per gallon (MPG), price, etc.
- Fit a $p = 2$ dim model. Plot m_i scores as points in \mathbb{R}^2 .
- To get a better understanding of the “meaning” of the latent factors, project unit vectors $e_1 = (1, 0, \dots, 0)$, $e_2 = (0, 1, 0, \dots, 0)$, etc. into the low dimensional space (blue lines)
- Horizontal axis represents price, corresponding to the features labeled “dealer” and “retail”, with expensive cars on the right. Vertical axis represents fuel efficiency (measured in terms of MPG) versus size: heavy vehicles are less efficient and are higher up, whereas light vehicles are more efficient and are lower down.
- Verify by finding the closest exemplars in the training set.

Example

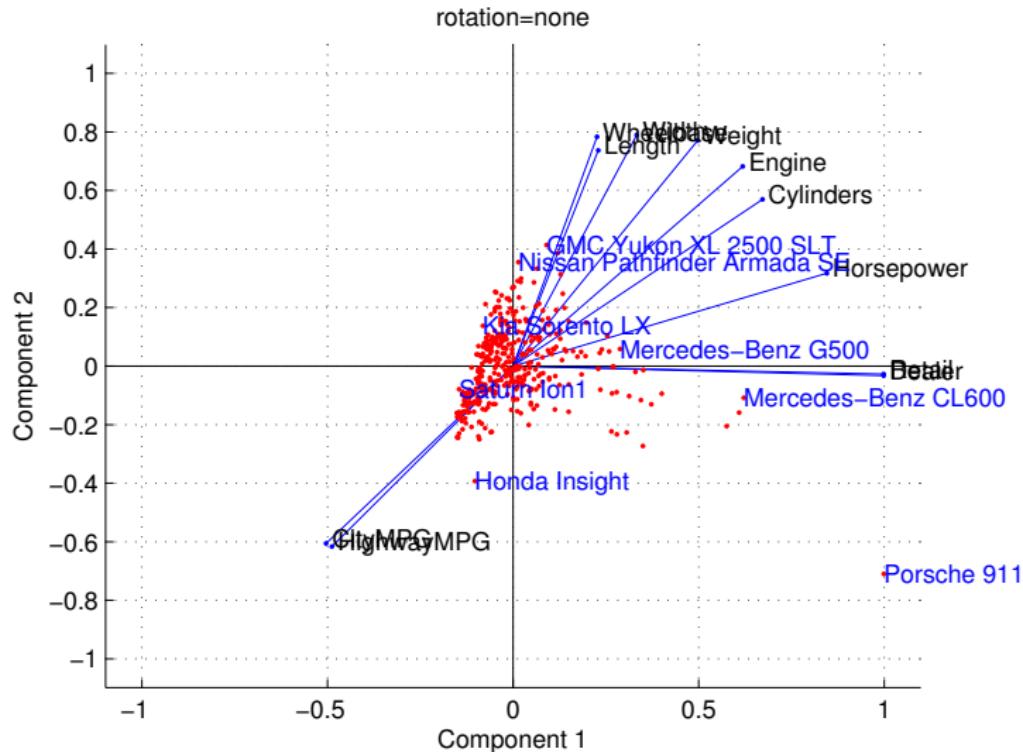
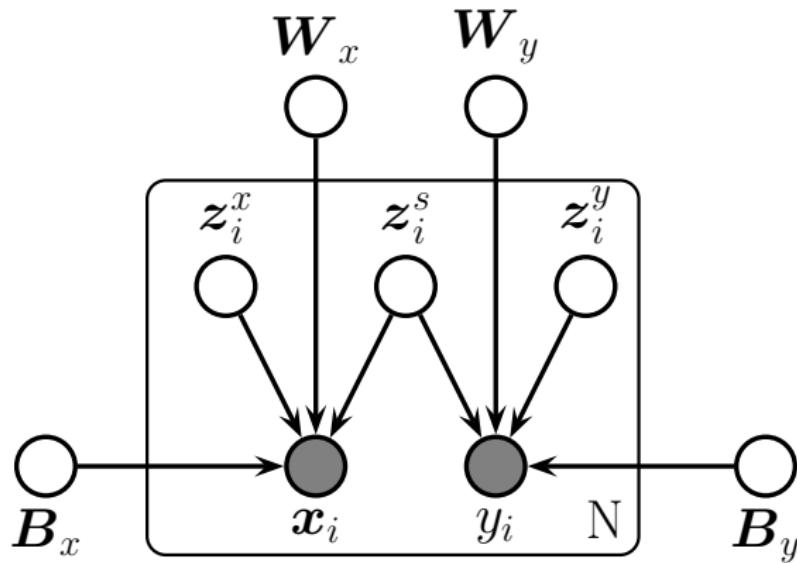


Figure 12.2 in K. Murphy

Special Cases: PCA and CCA

- Covariance matrix $\Psi = \sigma^2 I \rightsquigarrow$ (probabilistic) **PCA**.
- Two-view version involving x and $y \rightsquigarrow$ **CCA**.



From figure 12.19 in K. Murphy

PCA and dimensionality reduction

Given n data points in p dimensions:

$$X = \begin{bmatrix} - & \mathbf{x}_1 & - \\ - & \mathbf{x}_2 & - \\ - & \vdots & - \\ - & \mathbf{x}_n & - \end{bmatrix} \in \mathbb{R}^{n \times p}$$

Want to reduce dimensionality from p to k . Choose k directions $\mathbf{w}_1, \dots, \mathbf{w}_k$, arrange them as columns in matrix W :

$$W = [\mathbf{w}_1 \quad \mathbf{w}_2 \quad \dots \quad \mathbf{w}_k] \in \mathbb{R}^{p \times k}$$

For each \mathbf{w}_j , compute **similarity** $z_j = \mathbf{w}_j^t \mathbf{x}$, $j = 1 \dots k$.

Project \mathbf{x} down to $\mathbf{z} = (z_1, \dots, z_k)^t = W^t \mathbf{x}$. How to choose W ?

Encoding–decoding model

The projection matrix W serves two functions:

- **Encode:** $\mathbf{z} = W^t \mathbf{x}$, $\mathbf{z} \in \mathbb{R}^k$, $z_j = \mathbf{w}_j^t \mathbf{x}$.
 - ▶ The vectors \mathbf{w}_j form a basis of the projected space.
 - ▶ We will require that this basis is orthonormal, i.e. $W^t W = I$.
- **Decode:** $\tilde{\mathbf{x}} = W\mathbf{z} = \sum_{j=1}^k z_j \mathbf{w}_j$, $\tilde{\mathbf{x}} \in \mathbb{R}^p$.
 - ▶ If $k = p$, the above orthonormality condition implies $W^t = W^{-1}$, and encoding can be undone without loss of information.
 - ▶ If $k < p$, we use the pseudo-inverse
~~~ the reconstruction error will be nonzero.
- Above we assumed that the origin of the coordinate system is in the sample mean, i.e.  $\sum_i \mathbf{x}_i = 0$ .

# Principal Component Analysis (PCA)

In the general case, we want the reconstruction error  $\|\mathbf{x} - \tilde{\mathbf{x}}\|$  to be small.

Objective: minimize  $\min_{W \in \mathbb{R}^{p \times k}: W^t W = I} \sum_{i=1}^n \|\mathbf{x}_i - W W^t \mathbf{x}_i\|^2$

# Finding the principal components

Projection vectors are orthogonal  $\rightsquigarrow$  can treat them separately:

$$\begin{aligned} \min_{\mathbf{w}: \|\mathbf{w}\|=1} & \sum_{i=1}^n \|\mathbf{x}_i - \mathbf{w}\mathbf{w}^t\mathbf{x}_i\|^2 \\ \sum_i \|\mathbf{x}_i - \mathbf{w}\mathbf{w}^t\mathbf{x}_i\|^2 &= \sum_{i=1}^n [\mathbf{x}_i^t \mathbf{x}_i - 2\mathbf{x}_i^t \mathbf{w}\mathbf{w}^t\mathbf{x}_i + \mathbf{x}_i^t \mathbf{w} \underbrace{\mathbf{w}^t \mathbf{w}}_{=1} \mathbf{w}^t \mathbf{x}_i] \\ &= \sum_i [\mathbf{x}_i^t \mathbf{x}_i - \mathbf{x}_i^t \mathbf{w}\mathbf{w}^t\mathbf{x}_i] \\ &= \sum_i \mathbf{x}_i^t \mathbf{x}_i - \mathbf{w}^t \sum_{i=1}^n \mathbf{x}_i \mathbf{x}_i^t \mathbf{w} \\ &= \underbrace{\sum_i \mathbf{x}_i^t \mathbf{x}_i}_{\text{const.}} - \mathbf{w}^t \mathbf{X}^t \mathbf{X} \mathbf{w}. \end{aligned}$$

## Finding the principal components

- Want to maximize  $\mathbf{w}^t X^t X \mathbf{w}$  under the constraint  $\|\mathbf{w}\| = 1$
- Can also maximize the ratio  $J(\mathbf{w}) = \frac{\mathbf{w}^t X^t X \mathbf{w}}{\mathbf{w}^t \mathbf{w}}$ .
- Optimal projection  $u$  is the eigenvector of  $X^t X$  with largest eigenvalue (compare handout on spectral matrix norm).
- Note that we assumed that  $\sum_i \mathbf{x}_i = 0$ . Thus, the columns of  $X$  are assumed to sum to zero.
  - ~~ compute SVD of “centered” matrix  $X$
  - ~~ column vectors in  $W$  are eigenvectors of  $X^t X$
  - ~~ they are the principal components.

# Eigen-faces [Turk and Pentland, 1991]

- $p =$  number of pixels
- Each  $\mathbf{x}_i \in \mathbb{R}^p$  is a face image
- $x_{ji}$  = intensity of the  $j$ -th pixel in image  $i$

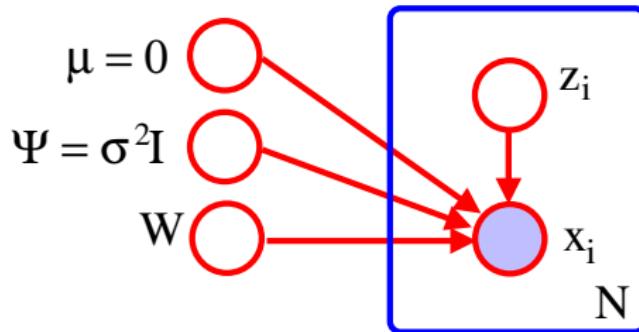
$$(X^t)_{p \times n} \approx W_{p \times k} (Z^t)_{k \times n}$$
$$\left( \begin{array}{c|c|c} \text{Face 1} & \dots & \text{Face n} \end{array} \right) \approx \left( \begin{array}{c|c|c|c|c} \text{Mean Face} & \text{Eigenface 1} & \text{Eigenface 2} & \text{Eigenface 3} & \text{Eigenface 4} \end{array} \right) \left[ \begin{array}{c|c|c} | & & | \\ z_1 & \dots & z_n \\ | & & | \end{array} \right]$$

Idea:  $\mathbf{z}_i$  more 'meaningful' representation of  $i$ -th face than  $\mathbf{x}_i$

Can use  $\mathbf{z}_i$  for nearest-neighbor classification

Much faster when  $p \gg k$ .

# Probabilistic PCA



- Assuming  $\Psi = \sigma^2 I$  and centered data in the FA model  
~~ likelihood

$$p(x_i|z_i, \theta) = \mathcal{N}(Wz_i, \sigma^2 I).$$

# Probabilistic PCA

- (Tipping & Bishop 1999): Maxima of the likelihood are given by

$$\hat{W} = V(\Lambda - \sigma^2 I)^{\frac{1}{2}} R,$$

where  $R$  is an arbitrary orthogonal matrix,

columns of  $V$ : first  $k$  eigenvectors of  $S = \frac{1}{n} X^t X$ ,

$\Lambda$ : diagonal matrix of eigenvalues.

- As  $\sigma^2 \rightarrow 0$ , we have  $\hat{W} \rightarrow V$ , **as in classical PCA** (for  $R = \Lambda^{-\frac{1}{2}}$ ).
- Projections  $\mathbf{z}_i$ : Posterior over the latent factors:

$$p(\mathbf{z}_i | \mathbf{x}_i, \hat{\theta}) = \mathcal{N}(\mathbf{z}_i | \hat{\mathbf{m}}_i, \sigma^2 \hat{F}^{-1})$$

$$\hat{F} = \sigma^2 I + \hat{W}^t \hat{W}$$

$$\hat{\mathbf{m}}_i = \hat{F}^{-1} \hat{W}^t \mathbf{x}_i$$

For  $\sigma^2 \rightarrow 0$ ,  $\mathbf{z}_i \rightarrow \hat{\mathbf{m}}_i$  and  $\hat{\mathbf{m}}_i \rightarrow V^t \mathbf{x}_i \rightsquigarrow$  orthogonal projection of the data onto the column space of  $V$ , **as in classical PCA**.